

▼ Predict survival on the Titanic

In this Lab, we ask you to apply the tools of machine learning to predict which passengers survived the tragedy

▼ Dataset

The dataset contains 891 observations of 12 variables:

- **PassengerId**: Unique ID for each passenger
- **Survived**: Survival (0 = No; 1 = Yes)
- **Pclass**: Passenger Class (1 = 1st; 2 = 2nd; 3 = 3rd)
- **Name**: Name
- **Sex**: Sex
- **Age**: Age
- **Sibsp**: Number of Siblings/Spouses Aboard
- **Parch**: Number of Parents/Children Aboard
- **Ticket**: Ticket Number
- **Fare**: Passenger Fare
- **Cabin**: Cabin
- **Embarked** Port of Embarkation (C = Cherbourg; Q = Queenstown; S = Southampton)

```
# import os
# from google.colab import drive
# drive.mount('/content/drive', force_remount=False)
```

```
# imports
import warnings
warnings.filterwarnings('ignore')
# your code here
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
titanic = pd.read_csv('titanic.csv',index_col=0)# your code here
titanic.head()
```

	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
PassengerId								
1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171
2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th	female	38.0	1	0	PC 17599

```
# print some info about the dataframe
# your code here
titanic.shape
```

```
(891, 11)
```

Looks like there are some Nan values, let's see how many for each column

```
titanic.isnull().sum()
```

```
Survived      0
Pclass        0
Name          0
Sex           0
Age          177
SibSp         0
Parch         0
Ticket        0
Fare          0
Cabin        687
Embarked      2
dtype: int64
```

Cabin contains a lot of Nan values, we'll drop this column

We'll replace the Nan values in **Age** with the age's median, and the ones in **Embarked** with 'S', which is the most frequent one in this column


```
# your code here to drop Cabin
titanic.drop('Cabin',axis=1,inplace=True)
# check the fillna documentation: http://pandas.pydata.org/pandas-docs/stable/ge
titanic["Age"] = titanic["Age"].fillna(value=titanic["Age"].median())
titanic["Embarked"] = titanic["Embarked"].fillna(value="S")
titanic.isnull().sum()
```

```
Survived      0
Pclass        0
Name          0
Sex           0
Age           0
SibSp         0
Parch         0
Ticket        0
Fare          0
Embarked      0
dtype: int64
```

▼ Visualization

```
%matplotlib inline
import matplotlib.pyplot as plt
print ('survival rate =', titanic.Survived.mean())
(titanic.groupby('Survived').size()/titanic.shape[0]).plot(kind="bar",title="Lab
```

```
survival rate = 0.3838383838383838
<AxesSubplot:title={'center':"Label's Balance"}, xlabel='Survived'>
```




```
# make a function to plot survival against passenger attribute
def survival_rate(column,t):
    df=pd.DataFrame()
    df['total']=titanic.groupby(column).size()
    df['survived'] = titanic.groupby(column).sum()['Survived']
    df['percentage'] = round(df['survived']/df['total']*100,2)
    print(df)

    df['survived'].plot(kind=t)
    df['total'].plot(kind=t,alpha=0.5,title="Survivors per "+str(column))
    plt.show()


# Draw survival per Sex
survival_rate("Sex","barh")
```

	total	survived	percentage
Sex			
female	314	233	74.20
male	577	109	18.89



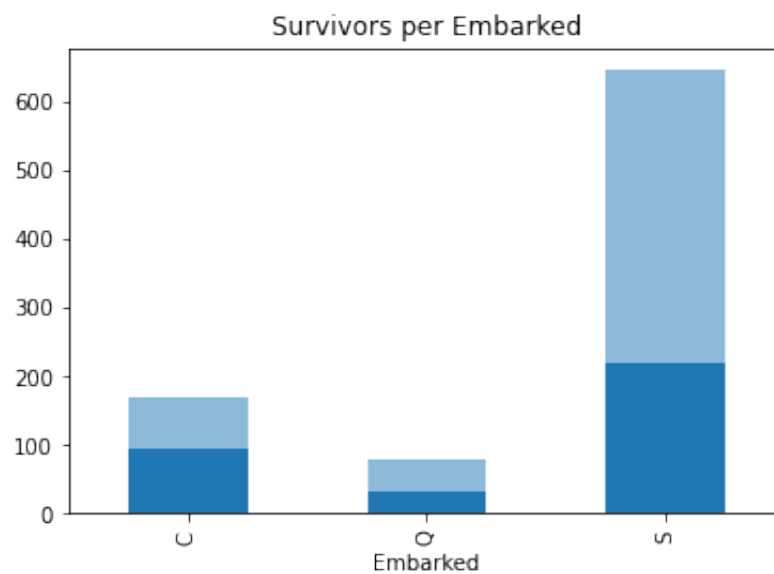
```
# Draw survival per Class
survival_rate("Pclass","barh")
```

	total	survived	percentage
Pclass			
1	216	136	62.96
2	184	87	47.28
3	491	119	24.24




```
# Graph survived per port of embarkation
survival_rate("Embarked","bar")
```

	total	survived	percentage
Embarked			
C	168	93	55.36
Q	77	30	38.96
S	646	219	33.90




```
# Draw survived per Number of Parents/Children Aboard (Parch)
# your code here
survival_rate("Parch","bar")
```

	total	survived	percentage
Parch			
0	678	233	34.37
1	118	65	55.08
2	80	40	50.00
3	5	3	60.00
4	4	0	0.00
5	5	1	20.00
6	1	0	0.00



```
# Draw survived per Number of Siblings/Spouses Aboard (SibSp)
# your code here
survival_rate("SibSp","bar")
```

	total	survived	percentage
SibSp			
0	608	210	34.54
1	209	112	53.59
2	28	13	46.43
3	16	4	25.00
4	18	3	16.67
5	5	0	0.00
8	7	0	0.00



▼ Model training

Some of the columns don't have predictive power, so let's specify which ones are included for prediction

```
predictors = ["Pclass", "Sex", "Age", 'SibSp' , 'Parch', "Fare", "Embarked"]
```

We need now to convert text columns in **predictors** to numerical ones

```

for col in predictors: # Loop through all columns in predictors
    if titanic[col].dtype == 'object': # check if column's type is object (text)
        titanic[col] = pd.Categorical(titanic[col]).codes # convert text to num

titanic.head()

```

	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	
PassengerId									
1	0	3	Braund, Mr. Owen Harris	1	22.0	1	0	A/5 21171	71.
2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th)	0	38.0	1	0	PC 17599	71.

```

# Split the data into a training set and a testing set. Set: test_size=0.3, random_state=42
# your code here
from sklearn.model_selection import train_test_split
y = titanic["Survived"]
x = titanic[predictors]
X_train, X_test, y_train, y_test = train_test_split(x,y,test_size=0.3, random_state=42)

print ("train shape", X_train.shape, y_train.shape)
print ("test shape", X_test.shape, y_test.shape)

```

```

train shape (623, 7) (623,)
test shape (268, 7) (268,)

```

```

# import LogisticRegression from: http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html
# your code here
from sklearn.linear_model import LogisticRegression
clf = LogisticRegression(random_state=1)
clf.fit(X_train,y_train)
# your code here
train_score = clf.score(X_train,y_train)
test_score = clf.score(X_test,y_test)
print ('train accuracy =', train_score)
print ('test accuracy =', test_score)

```

```

train accuracy = 0.8073836276083467
test accuracy = 0.7723880597014925

```


Let's print the model's parameters

```
coeff = pd.DataFrame()
coeff['Feature'] = X_train.columns
coeff['Coefficient Estimate'] = pd.Series(clf.coef_[0])
coeff.loc[len(coeff)] = ['Intercept', clf.intercept_[0]]
print (coeff)
```

	Feature	Coefficient Estimate
0	Pclass	-1.158692
1	Sex	-2.708761
2	Age	-0.040634
3	SibSp	-0.334012
4	Parch	0.071940
5	Fare	-0.000570
6	Embarked	-0.223307
7	Intercept	5.391544

We now need to predict class labels for the test set. We will also generate the class probabilities

```
# predict class labels for the test set
y_pred = clf.predict(X_test)
print (y_pred)
```

```
[1 0 1 1 1 0 0 1 1 1 0 1 0 0 1 0 0 0 0 1 0 0 1 0 1 0 1 1 0 1 1 0 0 1 0 1 0
 0 1 0 1 1 1 0 1 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 0 0 1 0 0 1 0 0 0 0 0
 1 0 1 0 0 1 0 0 0 0 1 0 0 0 1 1 0 0 0 1 0 1 0 1 0 0 1 0 0 1 1 0 0 0 0 0 0
 0 0 0 0 0 0 0 1 1 1 0 0 0 1 1 1 1 0 0 0 0 1 1 0 1 1 0 0 1 1 0 1 1 0 1 0 0
 1 0 1 0 0 1 0 0 0 0 1 0 0 0 1 0 0 1 1 0 0 0 1 1 1 0 1 0 0 0 1 0 1 1 0 0 1
 0 0 1 0 1 0 0 1 1 1 1 0 1 0 0 0 1 0 0 0 1 1 0 0 0 1 0 0 0 0 0 0 1 1 0 0 0
 0 0 0 0 1 0 1 0 0 0 0 0 0 1 1 1 1 0 0 0 0 0 0 0 1 1 0 1 0 0 0 1 1 1 0 1 0
 1 1 0 1 1 0 0 1 0]
```

```
# generate class probabilities : http://scikit-learn.org/stable/modules/generate
y_probs = clf.predict_proba(X_test)
print (y_probs)
```

```
[0.96819807 0.03180193]
[0.05269658 0.94730342]
[0.931658    0.068342   ]
[0.87604357 0.12395643]
[0.77746925 0.22253075]
[0.84945584 0.15054416]
[0.93503218 0.06496782]
[0.89643851 0.10356149]
[0.04125003 0.95874997]
[0.10522005 0.89477995]
```

```
[0.19533985 0.80466015]
[0.77085873 0.22914127]
[0.51933522 0.48066478]
[0.87382014 0.12617986]
[0.74132793 0.25867207]
[0.72964262 0.27035738]
[0.92394548 0.07605452]
[0.70094191 0.29905809]
[0.04415584 0.95584416]
[0.51470922 0.48529078]
[0.23425444 0.76574556]
[0.70097333 0.29902667]
[0.60811076 0.39188924]
[0.83497752 0.16502248]
[0.67771166 0.32228834]
[0.81308544 0.18691456]
[0.94205138 0.05794862]
[0.39282351 0.60717649]
[0.46814345 0.53185655]
[0.33317311 0.66682689]
[0.28618636 0.71381364]
[0.9221556  0.0778444 ]
[0.94286813 0.05713187]
[0.83457972 0.16542028]
[0.88251181 0.11748819]
[0.92906434 0.07093566]
[0.95096974 0.04903026]
[0.92442744 0.07557256]
[0.19105465 0.80894535]
[0.22468389 0.77531611]
[0.65007535 0.34992465]
[0.09251543 0.90748457]
[0.94833689 0.05166311]
[0.9001962  0.0998038 ]
[0.68228929 0.31771071]
[0.40806752 0.59193248]
[0.19561719 0.80438281]
[0.27749947 0.72250053]
[0.76966959 0.23033041]
[0.27027965 0.72972035]
[0.83450642 0.16549358]
[0.44297783 0.55702217]
[0.10338877 0.89661123]
[0.92394548 0.07605452]
[0.36586709 0.63413291]
[0.14332527 0.85667473]
[0.76730525 0.23269475]
[0.80676685 0.19323315]
[0.46946197 0.53053803]
[0.76143753 0.23856247]]
```

As you can see, the classifier outputs two probabilities for each row. It's predicting a 1 (Survived) any time the probability in the second column is greater than 0.5. Let's visualize it all together.

```
pred = pd.DataFrame({
    "Survived_original": y_test,
    "Survived_predicted": y_pred,
    "Survived_proba": np.transpose(y_probs)[1]
})
pred["Comparison"] = pred.Survived_original == pred.Survived_predicted
pred.head()
```

	Survived_original	Survived_predicted	Survived_proba	Comparison
PassengerId				
863	1	1	0.860689	T
224	0	0	0.084543	T
85	1	1	0.873336	T
681	0	1	0.634082	Fa
536	1	1	0.922157	T

▼ Confusion matrix

```
from sklearn import metrics
print (metrics.confusion_matrix(y_test, y_pred))
print (metrics.classification_report(y_test, y_pred))
```

```
[[129  24]
 [ 37  78]]
      precision    recall  f1-score   support

     0       0.78      0.84      0.81       153
     1       0.76      0.68      0.72       115

 accuracy          0.77       268
 macro avg       0.77      0.76      0.76       268
 weighted avg    0.77      0.77      0.77       268
```

As you can see, we can have the classification report for each class

▼ K-Fold Cross Validation

```
# import cross_validation from: http://scikit-learn.org/stable/modules/generated
# your code here
from sklearn.model_selection import cross_val_score
clf = LogisticRegression(random_state=1)
scores = cross_val_score(clf, titanic[predictors], titanic["Survived"], scoring=
## see model
print(scores)
# Take the mean of the scores (because we have one for each fold)
print(scores.mean())

[0.7877095  0.78651685 0.78089888 0.76966292 0.82022472]
0.7890025735986442
```

When you are improving a model, you want to make sur that you are really doing it and not just being lucky. This is why it's good to work with cross validation instead of one train/test split.

